



OPEN

Compute Project

AMD Motherboard Hardware v3.0

Authors: Bob Ogrey, Eugene Chung, Otto Joe

1 Scope

This document defines the technical specifications for the general purpose AMD motherboard that will be used in financial services Open Compute Project servers.

2 Contents

1	Scope	2
3	Overview.....	4
3.1	License.....	4
3.2	CAD Models	5
4	Product Descriptions and Requirements	6
5	Motherboard Features	7
5.1	Block Diagram for Base Design Implementation.....	10
5.2	Placement and Form Factor	11
5.3	CPU and Memory	12
5.4	Northbridge PCI-E Usage	14
5.5	Southbridge/Peripheral Bus Controller.....	14
6	Leveraged Design	15
6.1	HPC Server in 1U Chassis.....	16
6.2	General Purpose Motherboard in 2U Chassis	18
6.3	Storage Server Motherboard in 3U Chassis	20
7	Chassis Information	22
7.1	Supported Configurations	22
7.2	Front Panel Switches and Indicators	23
7.3	Heat Sinks	23
8.0	BIOS	23
8.1	BIOS Chip	23
8.2	BIOS Socket.....	23
8.3	BIOS Source Code	23
8.4	BIOS Power Optimization Features	23
8.5	BIOS Setup Menu	24
8.6	Console Redirect	25
8.7	PXE Boot	25
8.8	Other Boot Options	25

8.9	Remote BIOS Update.....	25
8.10	Event Log	26
9	System Management	27
9.1	Requirements.....	27
9.2	Management Port	29
9.3	Temperature sensors	29
9.4	Fan Connections	29
10.	Power System.....	29
10.1	Power Supply.....	29
10.2	Power connector Pin-out (TBD)	30
10.3	Support for Open Compute Infrastructure	30
10.4	Power Budget for the Open Platform board configurations.....	30
10.5	VRM Specifications.....	31
10.6	Power Sequencing.....	32
10.7	Power Optimized Hardware Features.....	32
11	I/O System	33
11.1	PCI-Express Slots.....	33
11.2	PCI-E Mezzanine Card.....	33
11.3	OPMA Management connector	33
11.3	Network.....	33
11.4	USB Interfaces	33
11.5	SATA	34
11.6	Debug Header.....	34
12	Environmental Requirements	35
12.1	Vibration and Shock	35
13.	Prescribed Materials	36
13.1	Disallowed Components.....	36
13.2	Capacitors and Inductors.....	36
14	Operating System Support	37
15	Adherence to AMD Motherboard Design Guide.....	37

3 Overview

This document describes the "AMD Financial Services Open Platform and System Solutions," hereafter referred to as the Open Platform board.

The AMD Financial Services Open Platform is a G34 motherboard that is optimized for cost and power, and designed to fit into a 1U, 1.5U, 2U, 3U or taller chassis to support different hard drive configurations. The chassis will house this 16" by 16.5" board and come in different U-heights depending on the use case. This system is intended as a low power offering with a feature set tailored to meet the needs of large data center operators and cloud computing service providers.

The AMD Financial Services Open Platform is intended to provide a universal, highly re-useable common motherboard that targets 70% to 80% of enterprise infrastructure of Wall Street Council members. Compatibility and deploy-ability are chief design objectives and boards should fit into both the Open Compute/Facebook datacenter infrastructure as well as traditional Enterprise 19" rack enclosures. This means the solution will provide a flexible power delivery system that supports both Facebook Open Compute & Traditional Enterprise Datacenter high-efficiency power supplies. Like previous Open Compute projects, this motherboard is power-optimized and barebones, designed to provide the lowest capital and operating costs. The Open Platform board enables customer driven (versus OEM driven) hardware core features and requirements and will drive an open standards approach that should generate volume and competitive pricing.

3.1 License

As of April 7, 2011, the following persons or entities have made this Specification available under the Open Web Foundation Final Specification Agreement (OWFa 1.0), which is available at <http://www.openwebfoundation.org/legal/the-owf-1-0-agreements/owfa-1-0>:

Facebook, Inc.

You can review the signed copies of the Open Web Foundation Agreement Version 1.0 for this Specification at <http://opencompute.org/licensing/>, which may also include additional parties to those listed above.

Your use of this Specification may be subject to other third party rights. THIS SPECIFICATION IS PROVIDED "AS IS." The contributors expressly disclaim any warranties (express, implied, or otherwise), including implied warranties of merchantability, non-infringement, fitness for a particular purpose, or title, related to the Specification. The entire risk as to implementing or otherwise using the Specification is assumed by the Specification implementer and user. IN NO EVENT WILL ANY PARTY BE LIABLE TO ANY OTHER PARTY FOR LOST PROFITS OR ANY FORM OF INDIRECT, SPECIAL, INCIDENTAL, OR CONSEQUENTIAL DAMAGES OF ANY CHARACTER FROM ANY CAUSES OF ACTION OF ANY KIND WITH RESPECT TO THIS SPECIFICATION OR ITS GOVERNING AGREEMENT, WHETHER BASED ON BREACH OF CONTRACT, TORT (INCLUDING NEGLIGENCE), OR OTHERWISE, AND WHETHER OR NOT THE OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

CONTRIBUTORS AND LICENSORS OF THIS DOCUMENT MAY HAVE MENTIONED CERTAIN TECHNOLOGIES THAT ARE MERELY REFERENCED WITHIN THIS DOCUMENT AND NOT LICENSED UNDER THE OWF CLA OR OWFA. THE FOLLOWING IS A NON-EXHAUSTIVE LIST OF MERELY REFERENCED TECHNOLOGY: PCI Express®, Hypertransport™, and USB. IMPLEMENTATION OF THESE TECHNOLOGIES MAY BE SUBJECT TO THEIR OWN LEGAL TERMS.

3.2 CAD Models

The following CAD files are incorporated by reference as if fully set forth in this specification:

TBD

4 Product Descriptions and Requirements

The following section describes the high level features for the various configurations of the Open Platform board. A single board design will be leveraged to meet the three usage cases listed below:

- 1) HPC server
- 2) General purpose server for virtualization and standard scale out
- 3) Storage Platform

However, in order to optimize for each usage case, a separate board sku with different component selection may be maintained. For example, the general purpose and HPC motherboard will populate only 1 of the 2 tunnel chips for IO expansion.

The common motherboard has the following features. It is a 16" x 16.5" board designed to fit into a 1U, 2U, or 3U chassis. The size of the motherboard was increased to support 3DIMMs per channel.

It has 2 processors, each with 12 memory sockets, 6 SATA connections per board, on-board 1Gb NIC with integrated management controller, up to four PCIe® expansion slots, mezzanine connector for custom module solutions, two serial ports, and two USB ports. Specific PCIe card support is dependent on usage case and chassis height.

System Configuration	Description (Common motherboard design is used for each of the following system configurations)
HPC	<ul style="list-style-type: none"> • 1 DIMM per channel (U/RDDR3 1600MHz, 1866 MHz* Stretch Goal) • Fits into the 1U chassis • Cooling and Power for SE 140W parts • 1 SR5670 tunnel • Supports 6 SATA drives natively off of the "Southbridge" • Supports up to ten 2.5" total drives with add-in card (Full details in section 7.2) • Supports up to 2 low-profile PCIe cards or 1

	standard height card <ul style="list-style-type: none"> • Single 1Gb on-board controller via BCM5725 • 10Gb solution via add-in mezzanine card
General Purpose	<ul style="list-style-type: none"> • 3 DIMM per channel (Up to 1600 MHz support) • Fits into 2U chassis • Cooling and Power for SE 140W Parts • Support for twenty five 2.5" SATA/SAS drives • Supports up to 1 standard-height and 1 low-profile PCIe cards (2 cards total) • Single 1Gb on-board controller via BCM5725 • 10Gb solution via add-in mezzanine card
Storage	<ul style="list-style-type: none"> • 3 DIMM per channel (Up to 1600 MHz support) • Fits into 3U chassis • Cooling and Power for SE 140W Parts • Support for thirty five 2.5" SATA /SAS drives • Supports up to 4 full-height, short PCIe cards • 10Gb solution via add-in mezzanine card

5 Motherboard Features

The base motherboard design supports the following feature set and will have different stuff options and implementations depending on the use case. A general description of the base motherboard design is documented below.

Component	Function
Socket	Socket G34
Processor	<ul style="list-style-type: none"> Two sockets per board Support for AMD processor codenamed "Magny-Cours", "Interlagos", and "Abu Dhabi" processors Abu Dhabi ("Orochi"-Rev C) support is mandatory Supports Infrastructure Group A, B, C: 85W, 115W, and 140W TDPs Magny-Cours: 8/12 cores codenamed "Greyhound" for Hydra die (MCM) Interlagos: 12/16 cores codenamed "Bulldozer" for Orochi die (SCM) Abu Dhabi: 4/8/12/16 cores codenamed "Piledriver" (MCM) Coherent Links: Triple x16 HyperTransport3 link supporting speeds up to 6.4 GT/s with support for HT1 operation @ 2.0 GT/s
Memory:	<ul style="list-style-type: none"> HPC config: 1 DIMM's per channel per processor (total of 4 DIMMs per processor and 8 DIMMs per system) Support up to 1866 MHz for 1 DIMM per channel General Purpose config: 3 DIMM's per channel per processor (total of 12 DIMMs per processor and 24 DIMMs per system) Three RDIMMs/channel up to 1600 MT/s 1DIMM/channel = 1600 MHz support 2DIMM/channel = 1333 MHz support 3DIMM/channel = 1066 MHz support Supports DDR3 1.5V, DDR3L (1.35V), DDR3U (1.25 V) Targeted configuration is 64GB, 1.35V, RDDR3, 1333 MHz Future LRDIMM (Load Reduced DIMM) support with AMD processor codenamed "Orochi" required UDIMM
I/O:	<ul style="list-style-type: none"> Two Tunnel SR56X0 Chipsets Single SR5670 for HPC and General Purpose Server One SR5670 and one SR5690 for Storage server Interconnect via nCHT3 16x16 link to North-Bridge. nCHT1 operation will not be supported with Orochi based products.

	<ul style="list-style-type: none"> ▪ AMD SP5100 South Bridge ▪ 6 onboard SATA 2.0 ports. ▪ 2 serial port ▪ 2 USB ports
Drive Bays	Support for 2.5" and 3.5" drives. Number of drives dependent on configuration and use case.
Video	No embedded support.
I/O Slots	<p>Up to 4 PCI Express® slots</p> <ul style="list-style-type: none"> ✓ Motherboard can support up to two x8 PCIe slots and two x16 PCIe slots (4 total) ✓ Three x8 PCIe slot support in a 1U or 2U chassis
Networking	Embedded BCM5725. 10GB via add-in mezzanine.
System Management	Integrated Management Controller via BCM5725.
System fans	Non-redundant and non-hot-swap, support cooling of Infrastructure Group A.
Board Dimensions	16" x 16.5"
Expected Layer Count	12 layers
Power target per MB	See Chapter 10.

5.1 Block Diagram for Base Design Implementation

Figure 1 illustrates the functional block diagram for the storage node version of the Open Platform board.

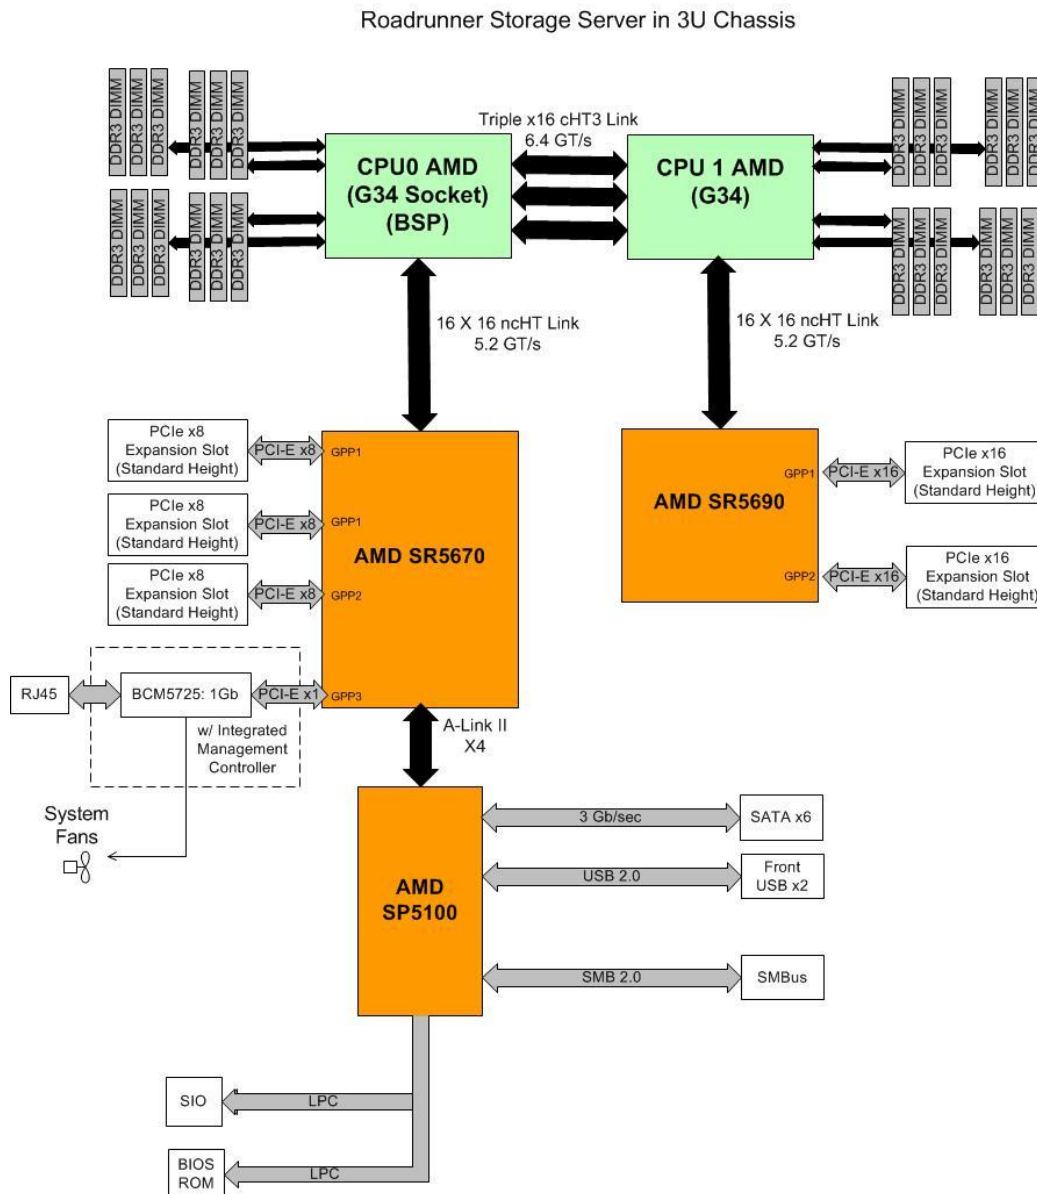


Figure 1 Functional Block Diagram

5.2 Placement and Form Factor

The motherboard's form factor is 16 x 16.5 inches with cutouts to support dual power supply locations. Figure 2 illustrates board placement. The placement shows the relative positions of key components, while exact dimension and position information will be determined in the future. Once released, the ODM should strictly follow the form factor, PCIe slot position, front IO port positions, PCIe mezzanine card connector position, power connector, and mounting holes, while other components can be shifted based on layout routing as long as relative position is maintained. Each specific configuration (HPC, general server, or storage) will use a slightly different flavor of the Open Platform board and chassis as shown in section 7.2.




Figure 2 AMD Motherboard Component Placement

5.3 CPU and Memory

The Socket G34 processor supports multi-chip modules (MCMs). The MCMs use HyperTransport™ technology as the inter-processor link. Figure 1 illustrates the internal HT link connections between the dies and the package.

Magny-Cours (6100 Series) uses the Greyhound core while Interlagos (6200 Series) uses the Bulldozer core. Abu-Dhabi is based on the Piledriver core will be available Q2 2012 and offers drop-in compatible part with 200 MHz performance uplift.

The Open Platform board should be designed to accommodate all three G34 power bands—85W, 115W, and 140 Interlagos and Abu Dhabi parts. The SE Abu Dhabi will consume up to 125A on the core and 25A on the north bridge. The board must support a 165A IDD spike and a 125A Load step on the core rail.

Interlagos in the G34 package has the following characteristic:

- 32 nm AMD Family 15h processor
- 16, 12 or 8 CPU cores - Each pair of cores comprises a “Bulldozer” module, a.k.a Compute Unit
- Each individual core has its own 16-Kbyte L1 data cache
- Each Bulldozer module has a single 64-Kbyte L1 instruction cache
- 2MB L2 cache per Bulldozer module
- Quad-channel U/RDDR-3 memory (also supports low voltage DDR3)
- Four HyperTransport™ technology links supporting speeds up to 6.4 GT/s.
- AMD-Virtualization™ (AMD-V™) technology with Rapid
- AMD Turbo CORE technology

5.3.1 DDR3

It is a requirement to support both 1.5V, 1.35 V and 1.25V memory offerings. This requires SPD detection by the system bios to detect the correct memory voltages and hardware interfaces to program the VDDIO voltage rail based on the populated memory configuration.

Future LR-DIMM support is a requirement and can be implemented with no adverse effects to standard DDR3 operation by following the chip select routing details. This involves routing the MA3_CS_L[1:0] to all DDR3 DIMM sockets in parallel.

Through-hole DIMM connectors are preferred due to cost, but a thorough analysis should be done to determine if SMT connectors provide improved memory margin or N+1 speed support.

4 channels DDR3 registered memory interface on each CPU:

- 3 DDR3 slots per channel per processor (total of 24 DIMMs on the motherboard)
- RDIMM/LV-RDIMM (1.5V/1.35V/1.25V), LRDIMM, and UDIMM/LV-UDIMM (1.35V/1.25V)
- SR, DR, and QR DIMMs
- DDR3 speeds of 800/1066/1333/1600/1866
- Up to maximum 768GB memory with 32GB RDIMMs

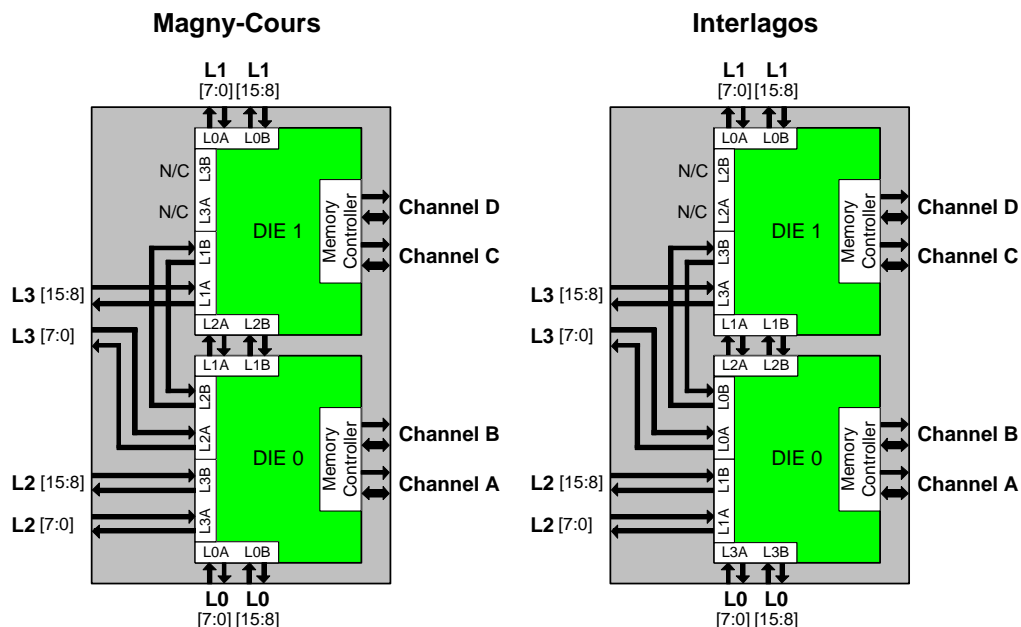


Figure 3: MCM Internal HyperTransport™ Interconnects

Figure 4 illustrates the internal memory bus connections between the dies and package.

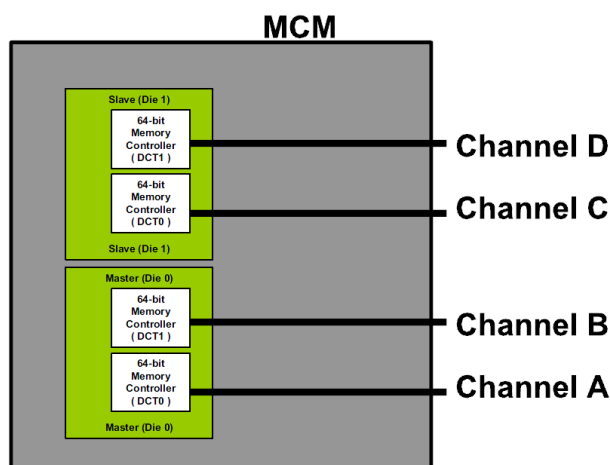


Figure 4: MCM Internal DDR Bus Connections

5.4 Northbridge PCIe Usage

The motherboard supports two SR5690 chipsets to provide maximum I/O in a balanced configuration. There are three different versions of the tunnel chipset, each with a different PCI-Express port configuration that is documented below.

	GPP1 Port	GPP2 Port	GPP3 Port	Total Available Lanes
SR5650	16 lanes	0 lanes	6 lanes	22 lanes
SR5670	16 lanes	8 lanes	6 lanes	30 lanes
SR5690	16 lanes	16 lanes	6 lanes	42 lanes

Depending on the usage case, the board will have different chipset population options. Please refer to section 6.0 for further details. For the storage configuration, a 5650 and a 5690 will be populated on the motherboard.

5.5 Southbridge/Peripheral Bus Controller

The motherboard uses the AMD SP5100 Southbridge chipset, which supports the following features:

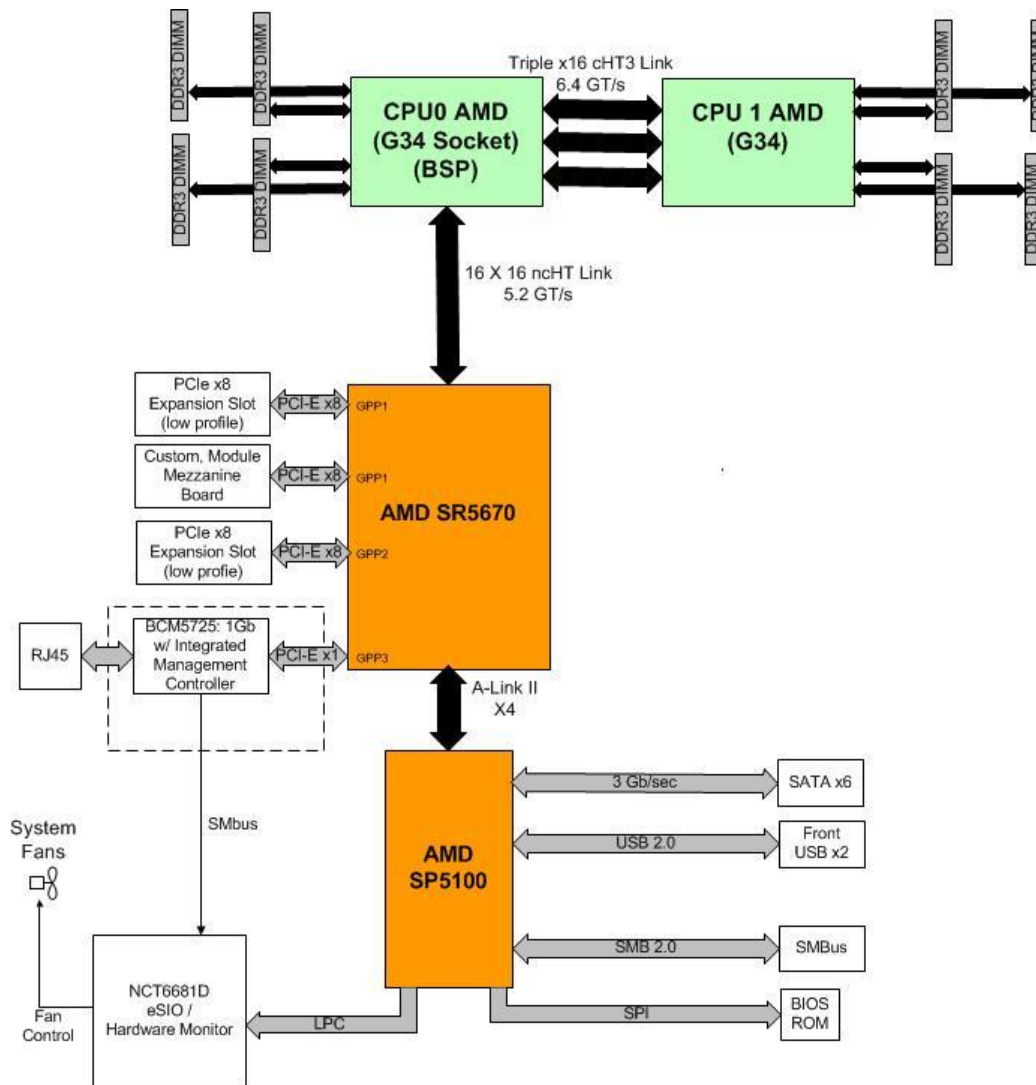
- 2 USB 2.0 ports (on the front panel)
- SATAII ports
- SPI interface
- SMBUS interface (master and slave)

6 Leveraged Design

The motherboard has been designed to be easily leveraged to support the HPC, general purpose, and storage configurations. All of the three use cases will be built around the same core design, but a few changes in both component population and the PCB will be made to optimize the design for the given application. The Base Model PCB will support two tunnel chips and 3DIMM per channel. Changes for the other configurations are listed below.

Segment	PCB	Component Change
HPC	Some memory traces are removed from the base model to support fastest possible 1DIMM per channel implementation.	Second I/O tunnel (56X0) is not populated.
General Purpose	Base Model PCB	Second I/O tunnel (56X0) is not populated.
Storage	Base Model PCB	SR5690 is populated to support two additional x16 PCIe slots. A SR5650 (in lieu of the 5690) can be populated if x8 support is sufficient.

6.1 HPC Server in 1U Chassis






Figure 5: HPC Motherboard Concept


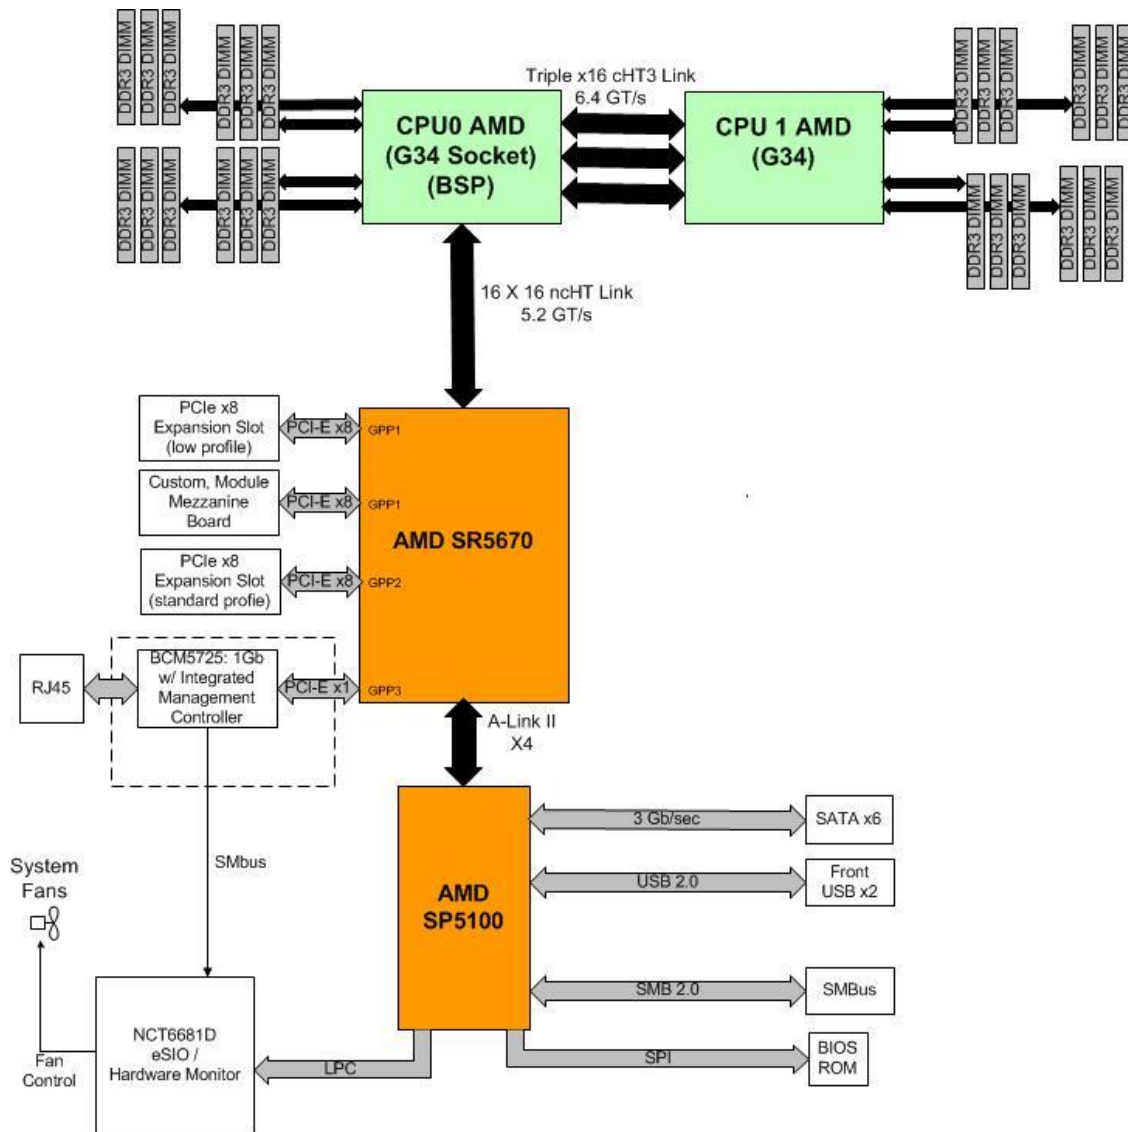


Figure 6: 1U HPC System Configuration 1

6.2 General Purpose Motherboard in 2U Chassis






Figure 8: General Server Motherboard


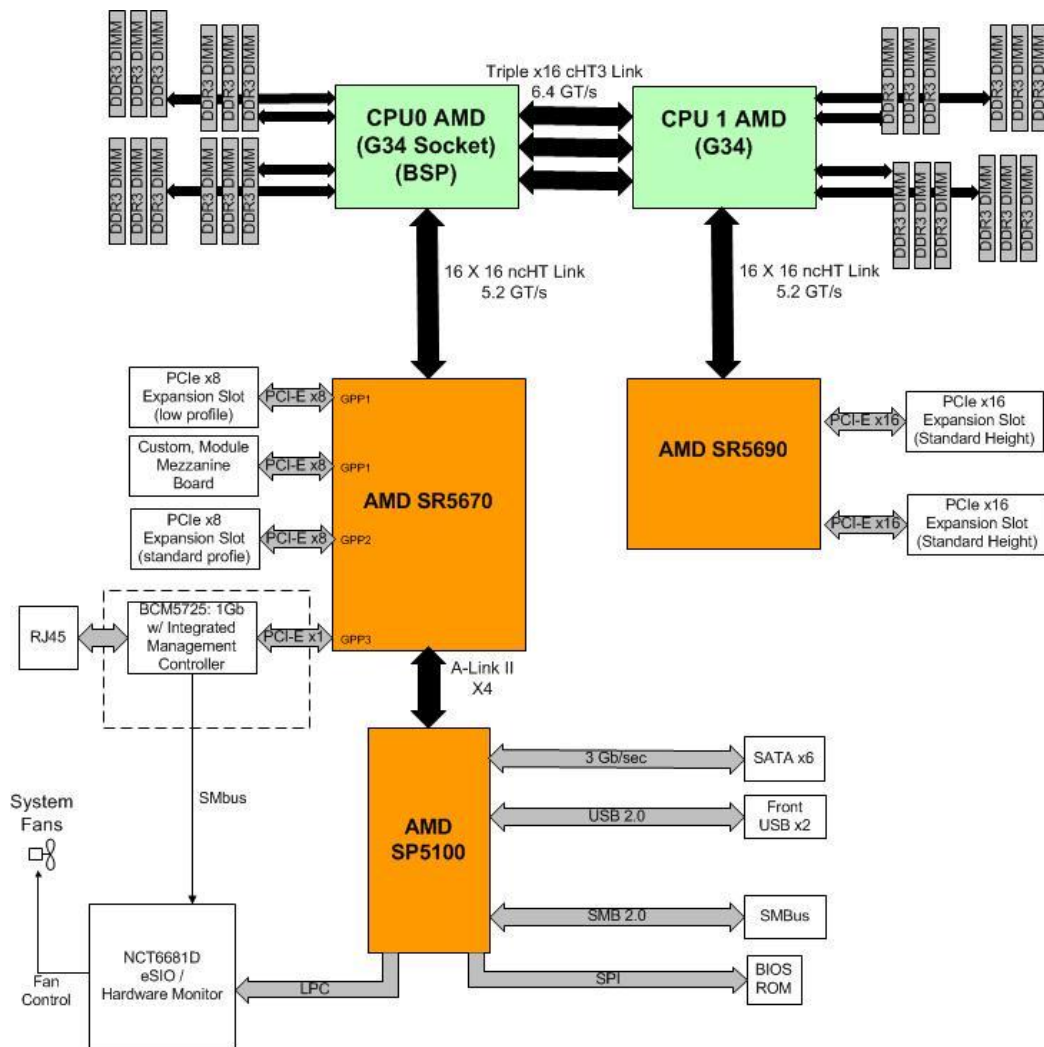


Fig 9: General Server in 2U Chassis 1

6.3 Storage Server Motherboard in 3U Chassis






Figure 10: Storage Server Motherboard




Fig 11: 3U Storage Server System

7 Chassis Information

The Open Platform board motherboard platform is designed to fit in a custom 1U/1.5U/2U/3U chassis.

The Open Platform board can fit into an Open Rack compliant chassis that will need to be developed. It can fit into future Open Compute1.0 chassis that have been designed with multiple holes punched into the base. However, it cannot fit into currently deployed Open Compute chassis because of incompatible hole patterns.

The 1U chassis is 730mm x 438.3mm x 43.6 mm and must be compatible with standard 19" EIA rack configurations.

The chassis is designed so that it can be easily assembled and the major components of the system are readily accessible.

7.1 Supported Configurations

The chassis will be available in several different heights but will maintain the same drive cage support mounting for the various U-heights. A support matrix is shown below.

Financial Services Matrix				
	1U	1.5U	2U	3U
Dual G34 socket	Up to 140W SE class	Up to 140W SE class	Up to 140W SE class	Up to 140W SE class
DIMMs	Up to 24 at 3DPC*	Up to 24 at 3DPC*	Up to 24 at 3DPC*	Up to 24 at 3DPC*
LP PCIe	2**	1**	1**	
Standard short PCIe		1**	1**	4***
3.5"/2.5" HDD	4/10	4/4	8/25	12/35
System fans	40x40x28 qty 12	60x60x38 qty 5	80x80x38 qty 8	120x120x38 qty 6
Management	Via BCM 5725 or add in module	Via BCM 5725 or add in module	Via BCM 5725 or add in module	Via BCM 5725 or add in module
Power	Redundant hot swappable multi voltage power supplies	Power adapter board with clips to Open Rack bus bar interface	Redundant hot swappable multi voltage power supplies	Redundant hot swappable multi voltage power supplies
10GbE	Via optional Mezz card or PCIe expansion	Via optional Mezz card or PCIe expansion	Via optional Mezz card or PCIe expansion	Via optional Mezz card or PCIe expansion
SAS	through optional PCIe expansion	through optional PCIe expansion	through optional PCIe expansion	through optional PCIe expansion
SATA	6 onboard ports	6 onboard ports	6 onboard ports	6 onboard ports
* HPC SKU variant will have 1DPC for a total of 8 DIMMs				
** requires 1U or 1.5U riser board				
*** no riser boards required, but requires 2nd 56XX chipset to be populate				

7.2 Front Panel Switches and Indicators

TBD based on feedback.

7.3 Heat Sinks

The motherboard supports heat sinks that are mounted according to the AMD G34 heat sink specification. The mounting device employs a back-plate and receptacles for screw-down type heat sinks. The ODM must comply with all keep out zones defined by AMD.

8.0 BIOS

The ODM is responsible for supplying and customizing a BIOS for the motherboard. The specific BIOS requirements are outlined in this section.

8.1 BIOS Chip

The BIOS uses the SP5100's SPI interface. The ODM is responsible for selecting a specific BIOS chip that meets the required functionality.

8.2 BIOS Socket

A socket on the motherboard holds the BIOS chip, which allows for manual replacement of the BIOS chip. The BIOS socket is easily accessible; other components on the motherboard do not interfere with the insertion or removal of the BIOS chip.

8.3 BIOS Source Code

The BIOS source code comes from TBD. The ODM is responsible for maintaining the BIOS source code to make sure it has latest code release from TBD and AMD.

8.4 BIOS Power Optimization Features

The BIOS is tuned to minimize system power consumption and should expose all device configuration, device feature, and power saving options in BIOS, to provide maximum flexibility to tune the system to the workload and operational environment. It should enable the following features:

- Unused devices disabled, including PCIe lanes, PCI, USB ports, and SATA/SAS ports

- Tuning CPU/chipset settings to reach minimized power consumption and best performance
- SPECpower is used as guidance for ODM to validate BIOS tuning results

Best practices include the following: :

- Full C1E support.
- "Efficient single fan variant" system fans that are non-redundant and utilize an aggressive fan management algorithm
- Disable HT Assist (Probe Filter Option)
- BIOS F2 setting(s) for HT Link Frequency. Enable cHT1 operation by default
- BIOS F2 setting(s) for HT link width. Enable 8-bit coherent link by default.
- Bios option to disable IOMMU operation
- Bios disable option for the BCM5275 NIC ports. '
- PSI options enabled by default
- Power Save options enabled
- A "green button" bios option to load all power optimized features
- Bios should enable CC6 (per core power gating) features in the processor. No specific HW changes need to be made to support this feature.

8.5 BIOS Setup Menu

The ODM must provide a BIOS specification, which includes the complete BIOS, setup menu, and default settings. The setup menu allows its options to be configured before the operating system loads. The configuration options available through the boot menu include the following:

- Setting for power feature after AC failure; default is set to keep last state.
- Setting for console redirection. Selectable options to support select console redirection from local COM port or the BMC's virtual UART for SOL.
- Setting for altitude of server deployment location.
- Hardware health monitoring display.
- Setting for watchdog timer; default is enabled and timeout value is 15 minutes.

- Event log viewing and clearing.
- Setting for ECC error threshold, available settings are 1, 4, 10, and 1000.
- Display power on self test (POST) results during boot up.
- If a CMOS checksum error happens (for example, caused by a BIOS update), the BIOS loads the system default automatically after showing a text message in the console for 5 seconds and rebooting the system to apply the update without user input.
- Setting to disable all "wait for keyboard input to continue" features.

8.6 Console Redirect

The BIOS detects the presence of a video card in the x16 PCIe slot. If a video card is present, the BIOS directs its output to the video card. If no video card is present, the BIOS directs its output to the board-mounted RS-232 console output.

8.7 PXE Boot

The BIOS supports PXE boot. When PXE booting, the system first attempts to PXE boot from the first Ethernet interface (eth0). If a PXE boot on the first Ethernet interface fails, the BIOS attempts to PXE boot from the second Ethernet interface (eth1).

8.8 Other Boot Options

The BIOS also supports booting from SATA/SAS and USB interfaces. The BIOS provides the capability to select boot options.

8.9 Remote BIOS Update

The BIOS can be updated remotely under these scenarios:

- Scenario 1: Sample/Audit BIOS settings
 - Return current BIOS settings, or
 - Save/export BIOS settings in a human-readable form that can be restored/imported (as in scenario 2)
- Scenario 2: Update BIOS with pre-configured set of BIOS settings
 - Update/change multiple BIOS settings
 - Reboot

- Scenario 3: BIOS/firmware update with a new revision
 - Load new BIOS/firmware on machine and update, retaining current BIOS settings
 - Reboot

Additionally, the update tool(s) should have the following capabilities:

- Update from the operating system over the LAN – the OS standard is CentOS v5.2
- Can complete update with a single reboot (no PXE boot, no multiple reboots)
- BIOS update or BIOS setup option change take no more than 5 minutes to complete
- No user interaction (like prompts)
- Can be scripted and propagated to multiple machines

8.10 Event Log

The BIOS logs system events through the baseboard management controller (BMC).

8.10.1 Logged Errors

- **CPU/memory errors:** Both correctable ECC and uncorrectable ECC errors should be logged into event log. Error categories include DRAM, HyperTransport Link, and L3 Cache.
- **HyperTransport errors:** Any errors that have a status register should be logged into the event log. Fatal or non-fatal classification follows the chipset vendor's recommendation.
- **Internal parity errors:** All errors which have status register should be logged into the event log. Fatal, non-fatal, or correctable classification follows the chipset vendor's recommendation.
- **PCIe errors:** All errors which have status register should be logged into Event Log, including root complex, endpoint device and any switch upstream/downstream ports if available. Link disable on errors should also be logged. Fatal, non-fatal, or correctable classification follows the chipset vendor's recommendation.

- **POST errors:** All POST errors detected by the BIOS during POST should be logged into the event log.
- **Power errors:**
- **MHOT and PROCHOT errors:** MEMHOT events should be logged with event source information indicating whether the event was triggered by a DIMM or a DIMM's Voltage Regulator. PROCHOT events should be logged with event source information indicating whether the event was triggered by a CPU or the CPU's Voltage Regulator.

8.10.2 Error Threshold Settings

An error threshold setting must be enabled for both correctable and uncorrectable errors. Once the programmed threshold is reached, an event should be triggered and logged.

- **Memory Correctable ECC:** The threshold value is 1000. When the threshold is reached, the BIOS should log the event including DIMM location information and output DIMM location code through the Facebook debug card.
- **HyperTransport errors:** Follow the chipset vendor's suggestion.
- **PCIe errors:** Follow the chipset vendor's suggestion.

9 System Management

The following sections describe the system management features implemented on the Open Platform board platform.

9.1 Requirements

- Out-of-band Management is provided by the integrated BMC (baseboard management controller) in the DCM5725 NIC.
- A separate NCT6681D eSIO/Hardware Monitor assists the BMC to provide fan speed control and fan/temperature monitoring.
- The integrated BMC provides out-of-band remote management access and control via: 1) the GbE network port in the BCM5725 or via a standard NC-SI interface that can be connected to a separate discrete NIC.
- The integrated BMC provides an IPMI/DCMI event log. BIOS will store memory or boot errors in the event log. The BMC will store sensor threshold events in the event

log. The event log can be read in-band or out-of-band using IPMI/DCMI standard interfaces.

- The BMC provides the following **out-of-band** remote control features:
 - Support for IPMI and DCMI network interfaces
 - Power-on/off and hard reset power control (via IPMI/DCMI or SMASH-2.0/WSManagement protocols)
 - Serial text console redirect (via IPMI SOL protocol, telnet, or SSH protocols)
 - IPMI Event log (via IPMI/DCMI)
 - Temperature sensors and hardware inventory (via IPMI/DCMI)
- The BMC provides the following **in-band** interfaces for systems management agents running on the host operating system or BIOS
 - KCS host HW/SW interface compliant w/ DCMI and IPMI
 - IPMI Event log rd/wr (via IPMI/DCMI commands)
 - Sensor and temperature monitoring (via IPMI/DCMI commands)
 - BMC Configuration (via IPMI/DCMI commands)
- The BMC also provides the following **advanced** management features:
 - Power capping per DCMI specification
 - Boot to a remote network CD image
 - Active Directory Authentication via Kerberos for SMASH 2.0/WSManagement and SSH console redirect out-of-band protocols
 - Web browser GUI for basic management functions

Notes:

DCMI (Data Center Management Interface) is a standard targeted for Data Center management. It uses interfaces defined in IPMI as the underlying protocol. DCMI minimizes the number of optional interfaces, requires support for specific temperature sensors, and adds power capping control to IPMI.

IPMI (Intelligent Platform Management Interface) is a widely used legacy management interface for servers. It uses a simple UDP-based network protocol and security protocols unique to IPMI.

SMASH (System Management Architecture for System Hardware) is the latest DMTF-defined management protocol. SMASH 2.0 defines a web-service protocol stack that uses SOAP-formatted messages, WSManagement, HTTP/S, TCP protocol stack with well accepted transport security.

9.2 Management Port

The onboard RJ45 can be configured as a dedicated management port or a shared 1-port GbE port. An on-board NC-SI connector will be populated on Ajax to support standard sideband interface to any add-in LAN cards.

9.3 Temperature sensors

The motherboard should support these thermal sensors:

- Two to monitor temperatures for CPU0 and CPU1, retrieved through the CPU's temperature sensor interface (TSI)
- Inlet temperature, retrieved through the thermistor, and located in the front of the motherboard
- Outlet temperature, retrieved through the thermistor, and located in the rear of the motherboard

9.4 Fan Connections

The motherboard has fan tachometer and PWM connections to control the fans listed in section 7.1.

10. Power System

This chapter describes the power sub-system for the Open Platform board platform.

10.1 Power Supply

The silver box should be specified and designed to provide the highest efficiencies possible. The chassis will be designed such that it supports a redundant power supply.

In addition, the power supply should meet the following criteria:

- Power supply should have an minimum 90% efficiency
- Power supply input should be TBD
- Multi-output power supply output
- PMBus interface support
- Redundant and hot-swappable power supply support

10.2 Power connector Pin-out (TBD)

TBD when power supply family is finalized.

10.3 Support for Open Compute Infrastructure

The Open Compute Infrastructure supports two power bus bars that run along both sides of the back of an Open Compute Rack.

The Open Platform board will fit seamlessly into the open compute infrastructure by utilizing a pcb adapter (red board in the pictures below) that connects to the bus bar. In addition to interfacing with the open rack bus bars, this pcb will take the place of the silver box supplies and provides any needed intermediate power conversion, such as 3.3V generation from 12V.

A conceptual drawing is shown below of Ajax in the Open Compute high PUE power distribution system.




Fig 12: Open Compute Support for the Open Platform board

10.4 Power Budget for the Open Platform board configurations

The summarized power budget below is shown for three different configurations—HPC, General Purpose, and Storage. HPC has the highest performing, highest TDP

CPUs, limited memory and IO. The 1U configuration is also the hardest configuration to cool and requires the highest power fans.

The general purpose configuration has more IO, standard power parts, fully populated DIMM slots, and uses SATA hard drives. The storage node is configured with 3 SAS controllers and fully populated memory.

It is important to note that these are maximum values. During production operation, bios controls can turn off features to significantly reduce power consumption from idle to 100%.

It is expected that the 3U storage node will use a different power supply than the 1U and 2U configurations.

		pwr	qty	SE Power		pwr	qty	Standard Power		pwr	qty	Standard Power		
Processors		140	2	280.0		115	2	230.0		115	2	230		
Ajax + CPU+ Memory + IO				376.8				422.8				445.9		
2.5" SSD		6	1	6.0										
SATA HDD	2.5" SATA					5.0	20	100.0		5.0	0	0		
SAS HDD	2.5" SAS 15k RPM	16.3	0	0.0		16.3	0	0.0		16.3	35	570.5		
Fans*		13.2	12	158.4		4.6	8	36.5		2.76	6	148.8		
System power (Twin + Ajax) - HPC Node				541.2	General Purpose-2U				559.2	Storage Node-4U				1,165.2

Table 4. Power Budget for the Open Platform board

10.5 VRM Specifications

The VRM and motherboard should support PSI features that allow the VRM to run in its peak efficiency range, i.e. redundant phases on the VRM will be turned off when the processor is operating at a low power point. The CPU VRM solutions should be load line enabled and designed to fit the specified power loads only.

In addition, the voltage regulators with the following characteristics are recommended:

- Load line to minimize load step transient.
- Differential Sense for accurate voltage sensing.
- Support for high value, low ESR ceramic caps and low ripple noise.
- In general, the target for ripple from the regulator should be < 10% of the maximum allowable voltage range. This allows 90% for all other sources of noise.

10.5.1 Target power efficiencies

The following are the target power efficiencies for each given group:

- Core VR: > 90% across the full operating load
- DDR VR: > 90% across the full operating load
- Silverbox: > 90% depending on load

10.6 Power Sequencing

The Open Platform board should follow power sequencing requirements for the G34, SP5650, and SP5100. Specific sequencing requirements are outlined in available design guides.

10.7 Power Optimized Hardware Features

Section 5.2 and Section 5.3 describes the high level power optimized hardware and BIOS features needed for the Open Platform board.

10.7.1 Switching Regulators versus Linear Regulators

Because the Open Platform board is designed to maximize system power savings, careful consideration must be undertaken when designing power delivery to the various system components. There are cost and board space considerations to balance, but it is generally recommended that all voltage rails that consume more than 1W of power be implemented using more efficient switching solutions.

Specific space, cost, and efficiency tradeoffs should be carefully analyzed during the earliest stages of the design.

10.7.2 VDDR implementation

In the past, VDDR and VLDT often shared a single regulator since both rails were specified for 1.2V operation. However, for low power designs, VDDR and VLDT should implement two separate regulators. VLDT will continue to operate at 1.2V but G34 has been updated to support 1.05V VDDR operation.

11 I/O System

This section describes the motherboard's I/O features.

11.1 PCI-Express Slots

A board fully populated with both SR56X0 chipset can support two x8 slots and two x16 slot. The second tunnel chip is a stuff option for heavy I/O configurations. If populated with a 5690, the second tunnel provides enough lanes to feed two x16 slots. If populated with a 5650, two x8 slots are available. Section 7.1 describes the mechanical limitations (riser needed, low profile vs. standard, etc.) for each chassis height.

11.2 PCIe Mezzanine Card

The motherboard has one mezzanine card connector that is intended for custom and modular 10Gb solutions. The mezzanine card dimension is W35.5mm x D262mm and expected solutions to be evaluated include the Broadcom 57810, Mellanox Connect-X-2, and Solar Flare SFC9000. The Mellanox solution is an x8 Gen2 PCIe solution, while the Broadcom 57810 is a Gen3 PCIe solution. Solar Flare is an x8, Gen2 PCIe solution.

11.3 OPMA Management connector

An OPMA management connector will be populated on the board. This connector accepts off the shelf management cards that can provide basic server management in lieu of the BCM5725.

11.3 Network

The motherboard has one BCM 5275 LAN chip on board to support the RJ45 connector. The BIOS supports PXE boot on the RJ45 port.

Each RJ45 connector has two built-in LEDs. While facing the RJ45 connector, the left LED is green single color; solid on means the link is active, and blinking means activity. The right LED is green/yellow dual color; green means 100M link speed while yellow means 1000M link speed.

11.4 USB Interfaces

The motherboard has two external USB ports located and the BIOS should support the following USB devices:

- Keyboard and mouse
- USB flash drive (bootable)
- USB hard drive (bootable)
- USB optical drive (bootable)

11.5 SATA

The motherboard has SP5100 interfaces on board, which support up to six SATA ports.

11.6 Debug Header

The motherboard includes a debug header on the front of the motherboard to display POST codes (see **Error! Reference source not found.**).The debug header supports hot plugging.

The debug card has two 7-segment LED displays, one RS-232 serial connector, and one reset switch. The RS-232 serial port provides console redirection. The two 7-segment LED displays show BIOS POST code and DIMM error information. The reset switch triggers a system reset when pressed.

The connector for the debug header is a 14-pin, shrouded, vertical, 2mm pitch connector. 0 is an illustration of the headers. The debug card has a key to match with the notch to avoid pin shift when plugging in.




Figure 13. Debug Header

Pin (CKT)	Function
1	Low HEX character [0] LSB
2	Low HEX character [1]
3	Low HEX character [2]
4	Low HEX character [3] most significant bit
5	High HEX character [0] least significant bit
6	High HEX character [1]
7	High HEX character [2]
8	High HEX character [3] most significant bit
9	Serial transmit (motherboard transmit)
10	Serial receive (motherboard receive)
11	System reset
12	Serial console select (1=SOL; 0=local)
13	GND
14	VCC (+5VDC)

Figure 14. Debug Header Pin Definition

12 Environmental Requirements

The motherboard meets the following environmental requirements:

Gaseous Contamination: Severity Level G1 per ANSI/ISA 71.04-1985

Ambient operating temperature range: -5°C to +45°C.

Operating and storage relative humidity: 10% to 90% (non-condensing)

Storage temperature range: -40°C to +70°.

Transportation temperature range: -55°C to +85°C (short-term storage).

The full OCP system also meets these requirements. In addition, the full system has an operating altitude with no de-ratings of 1000m (3300 feet).

12.1 Vibration and Shock

The motherboard meets shock and vibration requirements according to the following IEC specifications: IEC78-2-(*) and IEC721-3-(*) Standard & Levels. The testing requirements are listed in 0.

	Operating	Non-Operating
Vibration	0.5g acceleration, 1.5mm amplitude, 5 to 500 Hz, 10 sweeps at 1 octave/minute for each of the three axes (one sweep is 5 to 500 to 5 Hz)	1g acceleration, 3mm amplitude, 5 to 500 Hz, 10 sweeps at 1 octave/minute for each of the three axes (one sweep is 5 to 500 to 5 Hz)
Shock	6g, half-sine 11mS, 5 shocks for each of the three axes	12g, half-sine 11mS, 10 shocks for each of the three axes

Figure 15. Vibration and Shock Requirements

13. Prescribed Materials

13.1 Disallowed Components

The following components are not used in the design of the motherboard:

Components disallowed by the European Union's Restriction of Hazardous Substances Directive (RoHS 6)

Trimmers and/or potentiometers

Dip switches

13.2 Capacitors and Inductors

The following limitations apply to the use of capacitors:

Only aluminum organic polymer capacitors made by high quality manufacturers are used; they must be rated 105°C.

All capacitors have a predicted life of at least 50,000 hours at 45°C inlet air temperature, under worst conditions.

Tantalum capacitors are forbidden.

SMT ceramic capacitors with case size > 1206 are forbidden (size 1206 are still allowed when installed far from the PCB edge and with a correct orientation that minimizes risks of cracks).

Ceramic material for SMT capacitors must be X7R or better material (COG or NP0 type should be used in critical portions of the motherboard).

Only SMT inductors may be used. The use of through hole inductors is disallowed.

14 Operating System Support

Most up to date information about OS support is listed below.

ASSUMES latest updates/patches are installed*		
Enabled Optimized to support some or all of "Bulldozer's" new features	Compatible Will boot and run but not take advantage of "Bulldozer's" new features outside of new instructions	Not Supported Will not run on "Bulldozer" platforms and/or will not be supported by OSV
Includes new instruction support: <ul style="list-style-type: none"> Linux kernel 2.6.37 + , 3.0 + Novell SLES 11 SP2 (includes Xen) RHEL 6.2 with KVM (with latest z-stream patches) Windows Server 2008 R2 SP1 (optional scheduler patch available) Windows Server 2012/Hyper-V (in development) Xen 4.1 + Ubuntu 11.04 (w/ KVM) VMware vSphere 5.0 <p>Versions in this category also include latest software advances</p>	Includes new instruction support: <ul style="list-style-type: none"> Linux kernel 2.6.32 – 2.6.36 Novell SLES 11 SP1 RHEL 6.1 Ubuntu 10.10 Does <u>not</u> support new instructions for either Bulldozer or Sandy Bridge: <ul style="list-style-type: none"> Hyper-V R1 Hyper-V R2, Hyper-V R2 SP1 Novell SLES 10 SP4 and higher RHEL 5.7 (included KVM) Solaris 10u9, 11 VMware vSphere 4.1u2 Windows Server 2003 R2 SP2 Windows Server 2008 R2 Windows Server 2008 SP2 Xen 3.4.2 <p><i>Will run but not necessarily provide performance uplift</i></p>	<ul style="list-style-type: none"> Linux kernel 2.6.31 or earlier Novell SLES 10 thru SP3 Novell SLES 11 RHEL 4.x RHEL 5.0 – 5.5 RHEL 5.6 (can run with patches but is not supported by Red Hat) RHEL 6.0 Solaris 10 – 10u8 VMware ESX 3.5 VMware ESX 4.0 – 4.1u1 Windows Server 2003 versions prior to R2 SP2

Figure 16: Supported Operating Systems

15 Adherence to AMD Motherboard Design

The expectation is that close partnership between AMD and a selected ODM will occur. This will include joint review of specification, schematic, stack-up, layout, thermal, and chassis design.